

Applied & generative AI

Xander Vandooren

November 25, 2024

Contents

1	Session 7:	1
1.1	Recommender systems:	1
1.1.1	Gevaren:	1
1.1.2	Type input data:	1
1.1.3	Basic concepts:	1
1.1.3.1	Item-content matrix:	2
1.1.3.2	User-rating Matrix (URM):	2
1.1.3.2.1	Explicit ratings: Large rating scales vs small rating scales:	2
1.1.3.2.2	Explicit ratings: even vs odd ratings:	3
1.1.3.2.3	Explicit rating: user bias:	3
1.1.3.2.4	Top popular item:	3
1.1.3.2.5	Best rated item:	3
1.1.3.2.6	Best rated item with shrink term: .	3
1.1.4	Preprocessing: global effects formula:	4
1.1.5	Evaluating recommender systems:	4
1.1.5.1	Quality indicators:	5
1.1.5.2	on-line & off-line evaluation:	5
1.1.5.2.1	on-line: direct user feedback: . . .	5
1.1.5.2.2	on-line: A/B testing	5
1.1.5.2.3	on-line: controlled experiments: .	6
1.1.5.2.4	Off-line evaluation 4 steps:	6
1.1.6	Content-based filtering:	7
1.1.6.1	Recommendations based on cosine similarity	7
1.1.6.1.1	Dot product:	7
1.1.6.1.2	Improving on cosine similarity: shrink- ing	7
1.1.6.1.3	Estimating a rating:	7
1.1.6.2	K-nearest neighbours (KNN):	8
1.1.6.3	TF-IDF	8
1.1.6.3.1	ICM improvements: non-binary at- tributes:	8
1.1.6.3.2	ICM improvements: attribute weights:	8
1.1.6.3.3	ICM improvements: TF-IDF: . . .	8

1.1.6.3.4	Pros and cons content-based filtering:	9
1.1.7	Collaborative filtering:	9
1.1.7.1	User-rating matrix:	9
1.1.8	Model based vs memory based systems:	10

1 Session 7:

1.1 Recommender systems:

Wat zijn recommender systems?

- Traditional media:
 - Broadcasting is time limited

Grafiek: met recommenders proberen ze dit recht te trekken. Voorbeeld van recommender systems zijn spotify, Amazon, Netflix.

1.1.1 Gevaren:

- Je kan vast zitten in je eigen bubbel.
- Op gegeven moment zelfde genre (bvb bij netflix altijd zelfde genre)
- Marketing

1.1.2 Type input data:

- Items data (bvb in films alles van karakteristieken: genre, acteurs,...)
- Gebruiker data (Expliciete feedback bvb duim omhoog, 4/5 sterren en impliciete data bvb kijk je een film uit of hoeveel afleveringen van serie kijk je uit etc. Impliciete data is vaak belangrijker dan expliciet)
- Interactie items (Hoe reageren we op iets.)

Context: bvb dag in de week of feestdag, geslacht, ...

1.1.3 Basic concepts:

- Algorithms
 - Non-personalized
 - Personalized

- * Knowledge-based filtering (KBF) (redelijk algemeen bvb gewoon als man die leeftijds groep is zal hij wss deze film graag zien)
- * Content-based filtering (CBF) (items vergelijken met elkaar)
- * Collaborative filtering (CF) (wss bekendste systeem)
 - User based
 - Item based
 - Matrix factorization
- * Context-aware (CARS)
- * Hybrids
 - Deep learning, Factorization machines...

1.1.3.1 Item-content matrix: attributen → strenght of connection

1.1.3.2 User-rating Matrix (URM):

- Explicit:
 - ratings, thumbs-up...
- Implicit:
 - Did a user finish the movie?
 - Did he watch the second part?
 - Did he fast-forward?

User-item interaction matrix Meestal is er maar 1% van uw matrix ingevuld (groter dan 0). Bij netflix is de URM density=0.02%.

1.1.3.2.1 Explicit ratings: Large rating scales vs small rating scales:

- Large rating scale:
 - large user effort
 - fewer ratings

- Small rating scale:
 - small user effort
 - more ratings

1.1.3.2.2 Explicit ratings: even vs odd ratings:

- even ratings scale:
 - user forced to express opinion
 - few ratings
- odd ratings scale:
 - trend in giving neutral rating
 - more ratings
 - user more comfortable

1.1.3.2.3 Explicit rating: user bias: bvb. enkel negatieve reviews omdat mensen enkel hun tijd nemen om iets te raten als het slecht is.

1.1.3.2.4 Top popular item: Grootste hoeveelheid ratings

1.1.3.2.5 Best rated item: is eigenlijk gewoon formule gemiddelde doen.

1.1.3.2.6 Best rated item with shrink term: Items met veel hoge ratings zouden beter moeten zijn dan items met 1 top rating. Hier moet je een Determince C bias toevoegen dit is wel beetje trial en error om uit te zoeken wat juiste C moet zijn.

1.1.4 Preprocessing: global effects formula:

- Item bias: user will think differently about a movie now than in 20 years. Due to the style, genre, how innovative it is/was..
- User bias: users can rate a movie differently based on nostalgia, because they rate higher on average, because it reminds them of a book or a special moment...

1. Average ratings:

- Average ratings for all items and users

2. Normalized rating:

- To be computed for each user u and item i , only on non-zero ratings

3. Item bias:

- N_i : Number of users who have rated item i to be computed for each item i

4. Recompute rating:

- To be computed for each user u and item i , only on non-zero ratings

5. User bias:

- N_u : number of items i rated by user u to be computed for each user u

Final formula: Global effects formula to be computed for each user u and item i , only on non-zero ratings

1.1.5 Evaluating recommender systems:

Kwaliteit is afhankelijk van dataset → algorithm → user interface

1.1.5.1 Quality indicators:

- Relevance
- Coverage
- Novelty
- Diversity
- Consistency
- Confidence
- Serendipity

1.1.5.2 on-line & off-line evaluation:

1.1.5.2.1 on-line: direct user feedback: Ask for feedback by a rating, questionnaire.

- 2 problems:
 - Not always reliable opinions
 - Amount of feedback should be large enough

1.1.5.2.2 on-line: A/B testing Monitor the behaviour:

- Set A: users with a recommendation
- Set B: users without recommendation

Will they behave differently? (Buy more products, spend more time on Netflix...)

- But:
 - Difficult to set up
 - Difficult to interpret (if results bad: lack of relevance or lack of diversity?)

1.1.5.2.3 on-line: controlled experiments: Set up a mock-up application for a group of users

- Test application for a while and gather feedback

But:

- Not a real application, not the same as using real users

1.1.5.2.4 Off-line evaluation 4 steps:

- Defining task
 - Rating prediction (Top-N recommendation)
- Dataset
 - URM
- Partitioning
 - Model: look for connections in URM
 - User Profile: what does user like in general?
 - R: recommendation
 - Metric: compare recommendation with true opinion
 - **Three types of data needed:**
 - * Model uses data to detect rules -> $\text{model} = f(X)$
 - * Estimated ratings come from the model and user profile -> $\text{estimated ratings} = g(\text{model}, Y)$
 - * Compare estimated ratings with true value -> $\text{estimated rating} \leftrightarrow Z$
 - Oppassen voor overfitting!
- Metrics:
 - Quality metrics:
 - * Error metrics:
 - MAE & MSE

- * Classification metrics :
 - True positive
 - False positive
 - False negative
 - True negative
- * Ranking metrics:
 - ROC curve: Weten dat er een Area under curve is (AUC).
 - mAP (mean average precision)

1.1.6 Content-based filtering:

compare items based on their attributes. A user that expressed a preference for an item is likely to like similar items.

1.1.6.1 Recommendations based on cosine similarity an element of the ICM is 1 if the item has an attribute else 0.

1.1.6.1.1 Dot product: many attributes in common → the two items are very similar

↪ Cosine similarity is the normalization of the dot product. If two items are similar their cosine will be large.

1.1.6.1.2 Improving on cosine similarity: shrinking Shrinking:

- Reduce similarity to take into account only most similar with large support
- C: Shrink term

1.1.6.1.3 Estimating a rating: Weighted average of previous given ratings:

- r_{ui} = the rating of user u on item i
- Choose item with the highest r to recommend
- Or if you want an item of the top-N
- Remove the denominator to save calculation

1.1.6.2 K-nearest neighbours (KNN): Similarity matrix problems:

- Very dense - not many 0-values: computational heavy
- A lot of small similarity values:
 - A lot of very similar small values
 - Hard to distinguish between
 - Noise & low quality

Solution: keep K most similar items

- K influences the quality of the recommendations:
- Small K: not enough data for a reliable estimation
- Big K: too much noise in the data

1.1.6.3 TF-IDF

1.1.6.3.1 ICM improvements: non-binary attributes:

1.1.6.3.2 ICM improvements: attribute weights:

1.1.6.3.3 ICM improvements: TF-IDF: Do we have to manually figure out the attribute weights and non-binary attribute values? No! We can use a technique to do this automatically.

- TF-IDF: Term Frequency - Inverse Document Frequency
- $TF\text{-}IDF = TF * IDF \rightarrow$ The weight of each attribute depends on its frequency
- Replace Item-Content Matrix by TF-IDF matrix
- TF: Term frequency:
 - Higher if attribute appears more for an item
- IDF: Inverse document frequency:

- Higher if attribute appears more over all items

TF-IDF: higher value to terms that are frequent for an item, but rare in the whole collection of items. One value per item - attribute pair.

1. Apply TF-IDF on the Item-Content matrix to obtain a TF-IDF matrix
2. Construct the similarity matrix by applying cosine similarity on the TF-IDF matrix. Also use a shrink term if necessary.
3. Apply K-Nearest Neighbours on the similarity matrix

1.1.6.3.4 Pros and cons content-based filtering:

- Pros:

–

1.1.7 Collaborative filtering:

Often better quality recommendations compared to other methods. Large number of different techniques: user-based and item-based are seen here. Ratings that a user has not explicitly given to an item, can be inferred because the rating is correlated across various users and items.

1.1.7.1 User-rating matrix: Matrix containing past interactions between users and items

- Explicit ratings: for example rating from 1 to 5 (0 if missing)
- Implicit ratings: for example 1 if there is interaction with item and 0 otherwise

Implicit information is often more interesting than explicit. Most simple algorithms choose between explicit and implicit.

- Search for similar users and recommend items they like:
 - If two users give similar ratings to several items, we assume that they share the same opinion.
 - Better to have a number between 0 and 1

1.1.8 Model based vs memory based systems:

Model based: when a new user is added, the similarity matrix does not change the model significantly (Item based). Memory based: when a new user is added, the similarity matrix changes (User based).